

Customer Segmentation Clustering based on Demographics and Behaviors

Spencer Hirsch

April 26, 2024

Abstract: Identifying similarities between customers can be a useful strategy for companies when creating marketing campaigns. By analyzing the demographics and the purchasing behaviors of customers, a company can gain better insight into its customer base. With valuable insight, there is potential to release marketing campaigns targeted towards these groups to increase sales. By employing unsupervised machine learning we can better identify the characteristics of similar groups to aid in such campaigns. This study aims to utilize customer demographics and purchasing behaviors to identify trends and group consumers together based on these characteristics. Through the use of three different unsupervised machine learning clustering methods, we can identify groups with shared characteristics and attempt to optimize these groups with a variety of preprocessing and feature reduction methods to increase the effectiveness of these models in an attempt to find the most effective model for this study.

Keywords – Machine learning; customer segmentation; clustering algorithm

1 Introduction

Customer segmentation can be a powerful technique used by retailers to gain a better understanding of their target markets. For a business to be successful it must be able to cater to its customer base and address their specific needs. In the modern world, machine learning techniques can be utilized to make it easier for companies to better understand their customers. Customer segmentation can be used to break a larger group of people into smaller groups, separating them based on their shared characteristics, such as income or how much an individual spends on wine. This information can be especially useful for marketing strategies as companies would be able to share special benefits with these groups of people to draw them in.

The data used for this exploration can be seen in Table 1, all data refers strictly to their demo-

graphic information or their shopping behaviors.[1] This dataset was not explicitly stated to be from any sort of market or company, however, it is fair to assume that the data was collected from some sort of grocery store. This seems to be the case for all features, the only feature that seems to be odd for a grocery store is the "Amount spent on gold products". Nonetheless, the other features appear to refer to a grocery store, therefore for this study that is what we will assume.

1.1 Related Work

Customer segmentation can serve as a useful tool in many disciplines, Zhao discusses a handful of potential uses when it comes to the application of customer segmentation. Whether it be for marketing purposes as previously mentioned but also for exploration purposes. Zhao brings up the possibility of utilizing this

tool to determine new markets to enter.[2] This could be an incredibly useful tool for a company looking to expand the scope of its business. If a company identifies trends in their data indicating that the market is starting to lean in a specific direction, they would be able to tailor their products or services to meet this demand. Zhao’s discussion of customer segmentation doesn’t necessarily focus on strictly unsupervised machine learning methods, the article discusses two methods, recency, frequency, and monetary (RFM) and K-Mean clustering algorithms.[2] Rather than performing an exploration into the methods the article provides a general overview of the field at the time of publishing.

Rather than exploring potential methods of customer segmentation, Holy and Sokol bring customer segmentation into the real world. Their study deals with the study and development of a clustering algorithm to be utilized by a major Czech drugstore chain.[3] Holy and Sokol utilize a variety of clustering methods paired with K-Means to better understand the patterns of customers, RFM, as discussed in Zhao,[2] purchased products structure (PPS), and store mission (SM).

Similarly to the study that I have conducted here, John, Shobayo, and Ogunleye aim to study a variety of unsupervised machine learning methods, along with Principle Component Analysis and RFM.[4] The goal of the study is similar to mine, as they are aiming to find the most effective method for customer segmentation. The authors of this paper explore, K-Means clustering, Gaussian mixture model (GMM), density-based spatial clustering of applications with noise (DBSCAN), agglomerative clustering, and balanced iterative reduction and clustering using hierarchies (BIRCH).[4] The study proved to perform well with the authors being able to achieve a Silhouette Score, a measure of consistency in the cluster, of 0.8, through the use of PCA and GMM.[4]

Just as in the previous study, Hicham and Karim utilize several unsupervised machine-learning models for customer segmentation. Their paper proposes a clustering ensemble method of, DBSCAN, K-Means, Mini Batch K-Means, and MeanShift. Through the use of their method, they were able to achieve a Silhouette Score of 0.73.[5] This method proved to be

useful, DBSCAN itself was only able to achieve 0.72 itself, so there was marginal gain with their technique.[5]

Turkmen also perform a study using a variety of methods in hopes of achieving greater results, in their paper they weigh K-means clustering, hierarchical clustering, DBSCAN, and RFM.[6] The results of this study were not as great as the previous study using an ensemble method, Turkmen was able to achieve a Silhouette Score of 0.6 using the K-Means clustering algorithm.[6]

We can see that there have been numerous efforts and resources dedicated to studying the use of a variety of clustering methods for customer segmentation. In the following study, we will see the reoccurrence of clustering methods, such as K-Means clustering, Hierarchical clustering, and DBSCAN clustering as well as the introduction of new preprocessing methods that were not used in the above studies.

2 Methods

Three methods of unsupervised machine learning clustering methods were used, K-Means Clustering, Hierarchical Clustering, and DBSCAN. In addition to these three methods, a variety of preprocessing methods and feature reduction methods were used to reduce the dimensionality of the data. The goal of course to find the best combination of methods that provide the most efficient model. Just as John, Shobayo, and Ogunleye[4], Hicham and Karim[5], and Turkmen[6], the primary metric used in determining the success of the models was the Silhouette Score. This value helps us determine how well the clustering method is constructing the clusters. A Silhouette Score closer to 1 signifies that the clusters are more closely associated, whereas a score close to -1 means that they are not.

Once accessing the dataset, we start by manually preprocessing the data, this includes removing some of the features in the dataset as well as modifying categorical features of the dataset. Once this was complete, preprocessing using a variety of algorithms was employed, as well as feature reduction algorithms. Once this was complete, the three clustering mod-

Feature Name	Feature Description
Id	Unique customer identifier.
Year_Birth	Birth year of customer.
Education	Highest level of education obtained by customer.
Marital_Status	Marital status of customer.
Income	Annual income of customer.
Kidhome	Number of young children in the home.
Teenhome	Number of teenagers in the home.
Dt_Customer	Date when customer first enrolled.
Recency	Last visit of customer.
MntWines	Amount spent on wines.
MntFruits	Amount spent on fruits.
MntMeatProducts	Amount spent on meats.
MntFishProducts	Amount spent on fish.
MntSweetProducts	Amount spent on sweets.
MntGoldProds	Amount spent on gold products.
NumDealsPurchases	Number of purchases made as part of a discount promotion.
NumWebPurchases	Number of purchases made through website.
NumCatalogPurchases	Number of purchases made through catalog
NumStorePurchases	Number of purchases made in store.
NumWebVisitsMonth	Number of times customer has visited the website.

Table 1: Original columns in dataset after removal of marketing campaign information.

els were employed and underwent hyper-parameter tuning to find the optimal number of clusters for the highest Silhouette Score.

2.1 Preprocessing

To prepare the data for this exploration, it was necessary to clean up the data to adhere to the task at hand. Due to the nature of the project, it was fitting that the marketing campaign was removed from the dataset, the initial dataset contained five marketing campaigns as well as response information from the customers. For clustering based on demographics and behaviors, this information was not necessarily important.

In addition to this feature removal, much of the existing data in the dataset was categorical, rather than leaving objects in the data frame, the data was converted to contain only numeric values, this process included 1-hot encoding all categorical data and splitting the DateTime object into three new fea-

tures, month, day, and year. Following the manual deletion and modification of the original dataset, the new dataset consisted of 33 features.

After manual feature reduction and cleaning were performed, additional preprocessing was done on the data. This included utilizing a variety of methods: Standard Scaler, Robust Scaler, Quantile Transformation, and Log Transformation. All of these methods were used in an attempt to increase the effectiveness of the clusters. We can see that some of the preprocessing methods formed stronger clusters than others. Robust Scaler, seen in Figure 2, was the worst preprocessing method used when it came to clustering. The method that had the strongest clustering was the preprocessing performed with Quantile Transformation, as seen in Figure 3.

In Figure 1, we can see well-forming clusters, that appear to spread out more in a pedal shape. This is not a necessarily poor way to cluster the data, however, this method was not as effective as clustering as some of the others.




Figure 1: Visual representation of data preprocessed using Standard Scaler.




Figure 2: Visual representation of data preprocessed using Robust Scaler.




Figure 3: Visual representation of data preprocessed using Quantile Transformation.

Figure 2 contains the cluster for the worst performing preprocessing method, creating seemingly one complete cluster rather than the other preprocessing methods which created smaller more effective clusters.

Figure 3 created what appears to be the most effective method of preprocessing for this dataset, Quantile Transformation. Creating much smaller and more compact clusters compared to the other methods.

Lastly, Log Transformation, seen in Figure 4, performed similarly to how the Standard Scaler preprocessing performed. It created smaller clusters within a larger cluster of data.

2.2 Feature Reduction

Feature reduction was a necessary addition to this project, with the preprocessed dataset containing 33 features. Two feature reduction methods were tested to determine which method would be most effective. PCA was the first method tested and the results were very poor. Which led to the discovery of the T-distributed Stochastic Neighbor Embedding (t-SNE) method. This method turned out to be far more effective for this dataset. Figure 5 depicts the t-SNE feature reduction with the n.clusters parameter set to 2.




Figure 4: Visual representation of data preprocessed using Log Transformation.

This was the most effective feature reduction method. Additional `n_cluster` values were tested, `n_clusters` set to 3 are shown in Figure 6. We can see that the clustering was far less effective using this method. For this reason, the only feature reduction used going forward is the t-SNE method with `n_clusters` set to 2, this is used across all methods discussed below.

2.3 Clustering

For this study, three unsupervised machine learning clustering algorithms were employed, K-Means, Hierarchical, and DBSCAN. As previously mentioned, the metric that determines the highest-performing method is the Silhouette Score.

2.3.1 K-Means

The first method tested was the K-Means algorithm, this was paired with all of the previously mentioned preprocessing methods, however, with K-Means none of these methods proved to be useful. K-Means clustering paired with t-SNE was sufficient in achieving the highest Silhouette Score for this method. K-Means underwent hyper-parameter tuning where I attempted to optimize the Silhouette Score by increas-




Figure 5: Visual representation of data with t-SNE feature reduction with 2 `n_clusters`.




Figure 6: Visual representation of data with t-SNE feature reduction with 3 `n_clusters`.




Figure 7: Visual representation of dataset clustered with K-Means method with t-SNE feature reduction applied.

ing the number of clusters, for this case 2 clusters through 10 clusters were tested. K-Means with six clusters ended up being the best-performing model. This achieved a Silhouette Score of 0.3843, in comparison to the other methods discussed later in this paper this was the worst-performing model. Figure 6, visualizes the results returned using K-Means clustering. We can see that there were fairly well-constructed clusters for 4 of them, however, we can see some outliers towards the right-hand side of the data. Two individual points seemingly construct their own clusters. Upon this realization, the assumption would have been that four clusters would perform better, this however, was not the case. A common theme shared among all results of K-Means was that these points created their own clusters, except for the 2-cluster model.

2.3.2 Hierarchical

Hierarchical clustering was another method that was used for this project, similarly to K-Means, t-SNE feature reduction was used to reduce the dimensionality of the data. However, in addition to t-SNE, the dataset was preprocessed using Quantile Transformation. All preprocessing methods were used, however, the highest Silhouette Score was returned using this method paired with t-SNE. Hyper-parameter tuning was performed on this model as well, testing a




Figure 8: Visual representation of dataset clustered with Hierarchical method with t-SNE feature reduction and Quantile Transformation applied.

range of clusters between 2 and 100 to find the optimal number with the highest Silhouette Score. The best performing model contained 21 clusters. This turned out to be the best-performing model, with a Silhouette Score of 0.6483. We can see the resulting cluster visualization in Figure 8, we can see how vastly different the results are in comparison to the K-Means Clustering algorithm. The clusters were much smaller and more concise which allowed for the higher Silhouette Score.

2.3.3 DBSCAN

Just as shown with Hierarchical Clustering, DBSCAN was used in part with t-SNE and Quantile Transformation preprocessing. Again, all preprocessing methods were used for this model, however this combination returned the highest Silhouette Score. DBSCAN was not as effective as Hierarchical Clustering, however, it was the second best-performing method with a Silhouette Score of 0.5785. The visualization of the clusters can be seen in Figure 9, the clustering appears to be very similar to that of Figure 8, this is because the preprocessing method and feature reduction method used in both DBSCAN and Hierarchical Clustering were the same. This means that the results were solely reliant on the clustering algorithms themselves. Just as with the two previ-




Figure 9: Visual representation of dataset clustered with DBSCAN method with t-SNE feature reduction and Quantile Transformation applied.

ous models, DBSCAN went through hyper-parameter tuning in an attempt to increase the effectiveness of the model. Hyper-parameter tuning for DBSCAN focused on two hyper-parameters, epsilon and minimum samples, a brute force approach was taken, searching for a combination that yielded the highest Silhouette Score. In the case shown in Figure 9, epsilon was 2 and minimum samples was 6.

3 Results

Hierarchical Clustering paired with t-SNE and Quantile Transformation was the most effective model for grouping customers. With these clusters now, we can view the related information for these groups, to identify patterns in this data to see what these groups are like. Tables 2, 3, and 4 refer to selected clusters from the Hierarchical clustering model. The cluster consists of 21 clusters, all looking similar to the ones presented. With these selected clusters distinct patterns can be identified. We can see that one cluster has a significantly higher income than the others and has fewer discounted purchases. In addition to these two attributes, we see an increase in the amount spent on wines. Information such as this is especially useful to companies who are interested in learning more about their customers and their spending habits. For

Hieracical Clustering: Cluster 1

Feature Name	Cluster Average
Annual Income	\$49,770.84
Amount Meat	\$44.22
Amount Wines	\$82.01
Discount Purchases	3
Store Purchases	4
Web Purchases	3

Table 2: Extracted average information from Cluster 1 of Hierachical Clustering paired with t-SNE feature reduction and Qunatile Transfomation.

Hieracical Clustering: Cluster 4

Feature Name	Cluster Average
Annual Income	\$56,041.42
Amount Meat	\$171.61
Amount Wines	\$394.30
Discount Purchases	2
Store Purchases	6
Web Purchases	4

Table 3: Extracted average information from Cluster 4 of Hierachical Clustering paired with t-SNE feature reduction and Qunatile Transfomation.

example, Cluster 1 from Table 2, has the highest number of discounted purchases. This could be evidence that additional coupons would draw these customers in more, therefore we could look further at their purchasing history to better cater coupons to these groups to bring in more of their business.

4 Conclusion and Future Work

This paper studied the effectiveness of a variety of preprocessing, feature reduction, and unsupervised machine learning methods in an attempt to find the most effective model for customer segmentation. Utilizing unsupervised machine learning for customer segmentation has proven itself to be a useful tool. Unfortunately, this study was unable to achieve as significant results as some mentioned in the Related Works section, such as Hicham and Karim’s study.[5] However, achieving a Silhouette Score of 0.65 was not

Hieracical Clustering: Cluster 19

Feature Name	Cluster Average
Annual Income	\$74,421.19
Amount Meat	\$472.14
Amount Wines	\$557.78
Discount Purchases	1
Store Purchases	8
Web Purchases	5

Table 4: Extracted average information from Cluster 19 of Hierarchical Clustering paired with t-SNE feature reduction and Quantile Transformation.

too far off. After reading Hicham and Karim’s work, it would be interesting to follow their approach in creating an ensemble method with my selected models to attempt to increase the Silhouette Score. Another interesting addition would be using this model on another customer segmentation dataset and observing its performance. Overall, the study proved to be effective at clustering customers based on their demographic information and purchasing behaviors.

References

- [1] V. Patel, *Customer segmentation/clustering*, Kaggle, Accessed on April 22, 2024. Available online: <https://www.kaggle.com/datasets/vishakhapat/customer-segmentation-clustering/data>.
- [2] J. Zhao, “Customer segmentation application based on k-means,” *Applied and Computational Engineering*, vol. 47, pp. 242–247, Mar. 2024. DOI: 10.54254/2755-2721/47/20241400.
- [3] O. Sokol and V. Holý, “The role of shopping mission in retail customer segmentation,” *International Journal of Market Research*, vol. 63, no. 4, pp. 454–470, 2020. DOI: 10.1177/1470785320921011.
- [4] J. John, O. Shobayo, and B. Ogunleye, “An exploration of clustering algorithms for customer segmentation in the uk retail market,” *Analytics*, vol. 2, pp. 809–823, Oct. 2023. DOI: 10.3390/analytics2040042.
- [5] N. Hicham and S. Karim, “Analysis of unsupervised machine learning techniques for an efficient customer segmentation using clustering ensemble and spectral clustering,” *International Journal of Advanced Computer Science and Applications*, vol. 13, no. 10, 2022. DOI: 10.14569/ijacsa.2022.0131016.
- [6] B. Turkmen, “Customer segmentation with machine learning for online retail industry,” *The European Journal of Social and Behavioural Sciences*, vol. 31, no. 2, pp. 111–136, 2022. DOI: 10.15405/ejsbs.316.